

SET – 1

2)a) Explain about Timestamp based protocol.

Timestamp-Based Protocol in DBMS:

Timestamp-based protocol is a concurrency control method that uses **timestamps** to manage the order of transactions. Each transaction is assigned a **unique timestamp** when it starts, which determines its order of execution.

Key Points:

- Ensures **serializability** by ordering transactions based on timestamps.
- Each data item maintains:
 - **Read Timestamp (RTS)** – latest time it was read.
 - **Write Timestamp (WTS)** – latest time it was written.
- A transaction **T** is allowed to:
 - **Read(X)** only if $TS(T) \geq WTS(X)$
 - **Write(X)** only if $TS(T) \geq RTS(X)$ and $TS(T) \geq WTS(X)$
- If rules are violated, the transaction is **rolled back**.

Advantage:

- No deadlocks (no locking mechanism used).

Disadvantage:

- More rollbacks compared to locking protocols.

2)b) Write about Transaction Properties.

Transaction:

A **transaction** is a sequence of database operations performed as a single logical unit of work.


1. Atomicity

- Ensures all operations in a transaction are completed or none.
 - If any operation fails, changes are rolled back.
 - Maintains database in a consistent state.
-

2. Consistency

- Preserves database rules (e.g., constraints).
 - Transforms data from one valid state to another.
 - No violation of integrity occurs after the transaction.
-

3. Isolation

- Transactions execute independently.
- Intermediate results are hidden from other transactions.
- Prevents data conflicts in concurrent execution. 

4. Durability

- Committed changes are saved permanently.
- Survives system crashes or power failures.
- Ensures data is stored in non-volatile memory.

5)a) Write about ISAM.

ISAM (Indexed Sequential Access Method) is a database storage technique optimized for fast reads via indexed and sequential access.

Structure:

- **Data:** Sorted sequentially by key in fixed blocks.
- **Indexes:** Primary index (maps keys to block addresses); optional secondary indexes.

Pros:

- Fast read access for static/semi-static data (e.g., archives, reports).
- Simple, predictable performance.


Cons:

- Slow writes due to overflow management.
- Fragmentation over time, needing maintenance.

Use Cases: Legacy systems, read-heavy environments with minimal data changes.

Modern Alternatives:

- **B+ Trees:** Auto-balancing, better for dynamic data.
- **VSAM:** Enhanced IBM version with improved space management.



5)B) Compare any two File Organizations.

No.	Aspect	Sequential File Organization	ISAM (Indexed Sequential Access Method)
1.	Storage Structure	Records are stored one after another in a fixed order, usually based on a key field.	Stores data sequentially and uses an index with pointers to access records.
2.	Access Type	Supports only sequential access; suitable for batch processing.	Supports both sequential and direct access using index pointers.
3.	Insertion & Deletion	Inserting or deleting requires rearranging the entire file, which is time-consuming.	Uses an overflow area for insertions, making it easier but requires periodic reorg.
4.	Performance	Efficient only for full scans; slow for searches and updates in large files.	Fast for reads, especially in read-heavy workloads; less efficient for frequent writes.
5.	Suitability	Best for static data and simple applications like log files or backups.	Suitable for static or read-mostly environments like report generation systems.

6) Write the problems related to Decomposition.

Problems Related to Decomposition in DBMS:

Decomposition is the process of breaking a relation into two or more sub-relations. While it helps remove redundancy and anomalies, it may also lead to the following problems:

1. Loss of Information (Lossy Decomposition)

- Some original data may be lost if the decomposition is not done properly.
- The original relation cannot be recovered through natural joins.

2. Dependency Preservation Problem

- All functional dependencies from the original relation may not be preserved in the decomposed relations.
- Makes constraint enforcement difficult.

3. Join Dependency Issues

- Rejoining decomposed relations may produce extra or missing tuples.
- Leads to incorrect query results.

4. Increased Query Complexity

- Queries may need to join multiple tables, increasing execution time and processing cost.

Set – 2

2)a) Explain ACID properties.

1. Atomicity

- Ensures all operations in a transaction are completed or none.
 - If any operation fails, changes are rolled back.
 - Maintains database in a consistent state.
-

2. Consistency

- Preserves database rules (e.g., constraints).
 - Transforms data from one valid state to another.
 - No violation of integrity occurs after the transaction.
-

3. Isolation

- Transactions execute independently.
 - Intermediate results are hidden from other transactions.
 - Prevents data conflicts in concurrent execution.
-

4. Durability

- Committed changes are saved permanently.
- Survives system crashes or power failures.
- Ensures data is stored in non-volatile memory.



2)b) Explain conflict serializability with an example?

Conflict Serializability in DBMS:

Conflict serializability ensures that a schedule of concurrent transactions is equivalent to some serial schedule, based on conflicting operations.

Conflicting Operations:

Two operations conflict if:

- 1. They belong to different transactions,**
- 2. They access the same data item,**
- 3. At least one of them is a write.**

Example:

Consider the following schedule:

T1	T2
Read(A)	
	Read(A)
Write(A)	
	Write(B)

Conflicts:

- T1.Write(A) conflicts with T2.Read(A) $\Rightarrow T2 \rightarrow T1$**
- T2.Write(B) has no conflict with T1 (they access different items)**

Precedence Graph:

- Nodes: T1, T2**
- Edge: $T2 \rightarrow T1$**

Since the graph has no cycle, the schedule is conflict serializable.

Conclusion:

The schedule is equivalent to the serial order: T2 followed by T1.

5)a) Explain about inserting node in B+trees

Inserting a Node in B+ Trees:

Insertion in a B+ tree maintains the sorted order and balance of the tree. B+ trees store data only in leaf nodes, while internal nodes hold keys for navigation.

Steps for Insertion:

1. Locate the Leaf Node

- Traverse the tree from the root to the appropriate leaf node where the key should be inserted.

2. Insert into Leaf Node

- If there is space, insert the key in sorted order.

3. Split if Overflow Occurs

- If the node overflows (i.e., exceeds maximum capacity):
 - Split the node into two.
 - Move the middle key to the parent node for redirection.

4. Repeat Splitting if Needed

- If the parent also overflows, repeat the split up to the root.
- If the root splits, a new root is created, increasing the tree height.

Example:

Suppose the order of B+ tree is 3 (max 2 keys per node):

Insert keys: 10, 20, 5, 6

- Insert 10 → fits in leaf
- Insert 20 → fits
- Insert 5 → causes overflow → split: [5], [10, 20] → promote 10
- Tree structure updates with 10 in root and two child leaves

This maintains balance, sorted order, and efficient access.

5)b) Write about ISAM.

ISAM (Indexed Sequential Access Method) is a database storage technique optimized for fast reads via indexed and sequential access.

Structure:

- **Data:** Sorted sequentially by key in fixed blocks.
- **Indexes:** Primary index (maps keys to block addresses); optional secondary indexes.

Pros:

- Fast read access for static/semi-static data (e.g., archives, reports).
- Simple, predictable performance.

Cons:

- Slow writes due to overflow management.
- Fragmentation over time, needing maintenance.

Use Cases: Legacy systems, read-heavy environments with minimal data changes.

Modern Alternatives:

- **B+ Trees:** Auto-balancing, better for dynamic data.
- **VSAM:** Enhanced IBM version with improved space management.

7) Describe about 1st Normal form?**First Normal Form (1NF):**

1NF is the basic level of normalization in relational databases. A relation is in 1NF if:

Rules of 1NF:**1. Atomic Values Only**

- Each cell must contain only a single (indivisible) value.
- No sets, arrays, or lists.

2. Unique Column Names

- Each column must have a unique name.

3. No Repeating Groups

- There should be no multiple columns for the same type of data.

Example (Before 1NF):

StudentID	Name	Courses
-----------	------	---------

1	Raju	DBMS, OS
2	Meena	DBMS

Here, "Courses" has multiple values → violates 1NF.

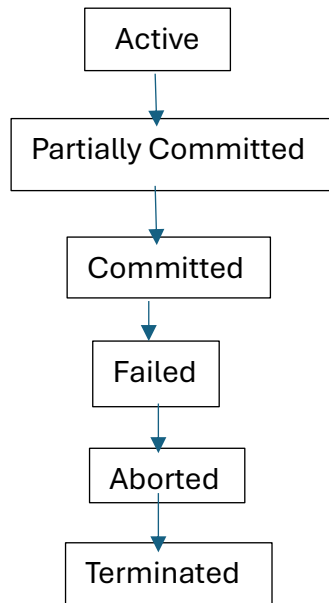
Converted to 1NF:

StudentID	Name	Course
1	Raju	DBMS
1	Raju	OS
2	Meena	DBMS

Conclusion:

1NF removes multi-valued attributes and ensures each field contains atomic data, making the table easier to manage.

2)a) Draw transaction state diagram and describe each state that a transaction goes through during its execution.



Transaction States Explained:

1. Active

- The transaction is currently executing its operations (read/write).

2. Partially Committed

- All operations are done, and the transaction is ready to save changes.

3. Committed

- Changes made by the transaction are permanently saved in the database.

4. Failed

- Some error occurred, so the transaction cannot proceed.

5. Aborted

- The transaction is rolled back, undoing all changes.

6. Terminated

- The transaction has either committed or aborted, and its execution is complete.

2)b) Explain ACID properties.

1. Atomicity

- Ensures all operations in a transaction are completed or none.
 - If any operation fails, changes are rolled back.
 - Maintains database in a consistent state.
-

2. Consistency

- Preserves database rules (e.g., constraints).
 - Transforms data from one valid state to another.
 - No violation of integrity occurs after the transaction.
-

3. Isolation

- Transactions execute independently.
 - Intermediate results are hidden from other transactions.
 - Prevents data conflicts in concurrent execution.
-

4. Durability

- Committed changes are saved permanently.
- Survives system crashes or power failures.
- Ensures data is stored in non-volatile memory.



4)a) Explain in detail about Data on External Storage in DBMS.

Data on External Storage in DBMS (Short Notes):

External storage refers to storing data on non-volatile devices like hard disks or SSDs, used when data is too large for main memory.

Key Concepts:

- **Block:** Unit of data transfer between disk and RAM.
- **File:** Contains multiple blocks, holds tables, indexes, etc.
- **Record/Tuple:** A row in a table, stored inside blocks.

Types of File Organization:

1. **Heap:** No order; fast insertions.
2. **Sequential:** Sorted by key; good for range queries.
3. **Hashed:** Uses hash function; fast for exact lookups.
4. **Indexed:** Uses indexes (e.g., B+ Trees) for fast search.

Uses:

- Stores large data permanently.
- Supports efficient retrieval and updates.
- Essential for big databases.

ISAM (Indexed Sequential Access Method) is a database storage technique optimized for fast reads via indexed and sequential access.

Structure:

- **Data:** Sorted sequentially by key in fixed blocks.
- **Indexes:** Primary index (maps keys to block addresses); optional secondary indexes.

Pros:

- Fast read access for static/semi-static data (e.g., archives, reports).
- Simple, predictable performance.

Cons:

- Slow writes due to overflow management.
- Fragmentation over time, needing maintenance.

Use Cases: Legacy systems, read-heavy environments with minimal data changes.

Modern Alternatives:

- **B+ Trees:** Auto-balancing, better for dynamic data.
- **VSAM:** Enhanced IBM version with improved space management.

7)Write about Decomposition & its problems in DBMS.

Problems Related to Decomposition in DBMS:

Decomposition is the process of breaking a relation into two or more sub-relations. While it helps remove redundancy and anomalies, it may also lead to the following problems:

1. Loss of Information (Lossy Decomposition)

- Some original data may be **lost** if the decomposition is not done properly.
- The original relation **cannot be recovered** through natural joins.

2. Dependency Preservation Problem

- All **functional dependencies** from the original relation **may not be preserved** in the decomposed relations.
- Makes **constraint enforcement** difficult.

3. Join Dependency Issues

- Rejoining decomposed relations may produce **extra or missing tuples**.
- Leads to **incorrect query results**.

4. Increased Query Complexity

- Queries may need to **join multiple tables**, increasing **execution time** and **processing cost**.